

Introduction to Artificial Intelligence: AO* and Hill Climbing algorithms

Dr. Tapas Kumar Mishra

Assistant Professor
Department of Computer Science and Engineering
SRM University
Amaravati-522502, Andhra Pradesh, India
Email: tapaskumar.m [at] srmap [dot] edu [dot] in



Heuristic Function

Admissible heuristics

E.g., for the 8-puzzle:

$h_1(n)$ = number of misplaced tiles

$h_2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$$h_1(S) = ?? \quad 6$$

$$h_2(S) = ?? \quad 4+0+3+3+1+0+2+1 = 14$$



Heuristic Function

- A typical solution is about **20 steps**, although this of course varies depending on the **initial state**.
- The **branching factor** is about **3** (when the **empty tile** is in the **middle**, there are **four possible moves**; when it is in a **corner** there are **two**; and when it is along an **edge** there are **three**).
- An exhaustive search to depth **20** would look at about $3^{20} = 3.5 \times 10^9$ states.
- There are only $9! = 362,880$ different arrangements of 9 squares.
- Manhattan/City block distance: $|x_1 - x_2| + |y_1 - y_2|$



Dominance

If $h_2(n) \geq h_1(n)$ for all n (both admissible)
then h_2 dominates h_1 and is better for search

Typical search costs:

$d = 14$ IDS = 3,473,941 nodes

$A^*(h_1) = 539$ nodes

$A^*(h_2) = 113$ nodes

$d = 24$ IDS \approx 54,000,000,000 nodes

$A^*(h_1) = 39,135$ nodes

$A^*(h_2) = 1,641$ nodes

Given any admissible heuristics h_a, h_b ,

$$h(n) = \max(h_a(n), h_b(n))$$

is also admissible and dominates h_a, h_b



Relaxed problems

Admissible heuristics can be derived from the **exact** solution cost of a **relaxed** version of the problem

If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution

If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

Key point: the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem



Local Search Algorithms and Optimization Problems

- **Local search algorithms** operate using a single **current state** (rather than **multiple paths**) and generally **move only to neighbors** of that state.
- Typically, the **paths** followed by the search are **not retained**.
- Although local search algorithms are **not systematic**, they have two key advantages:
 - 1 They use **very little memory**-usually a constant amount
 - 2 They can often **find reasonable solutions** in large or infinite (continuous) state spaces for which **systematic algorithms are unsuitable**.




Local Search Algorithms and Optimization Problems

- Local search algorithms are useful for solving pure **optimization problems**, in which the aim is to find the best state according to an **objective function**.
- A landscape has both “**location**” (defined by the **state**) and “**elevation**” (defined by the value of the **heuristic cost function** or **objective function**).
 - If elevation corresponds to **cost**, then the aim is to find the **lowest valley-a global minimum**
 - If elevation corresponds to an **objective function**, then the aim is to find the **highest peak-a global maximum**.



Hill-Climbing Search




- It is simply a loop that **continually** moves in the direction of **increasing value**-that is, **uphill**.
- It terminates when it reaches a “**peak**” where **no neighbor** has a **higher value**.
- Hill-climbing algorithms typically choose randomly among the set of best successors, if there is more than one.

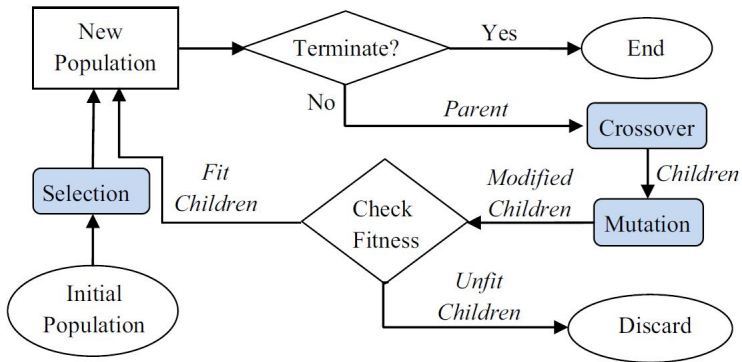


Hill-Climbing Search

Hill climbing often gets stuck for the following reasons


- 1 **Local maxima:** a **local maximum** is a **peak** that is **higher** than each of its **neighboring states**, but **lower** than the **global maximum**.
- 2 **Ridges:** Ridges result in a **sequence of local maxima** that is very difficult for **greedy algorithms** to **navigate**.
- 3 **Plateaux:** A plateau is an area of the state space landscape where the **evaluation function is flat**. It can be a **flat local maximum**, from which no uphill exit exists, or a **shoulder**, from which it is possible to make progress.






Genetic algorithms


= stochastic local beam search + generate successors from **pairs** of states



AND OR Search



AND OR Search





AND/OR graphs


- Some problems are best represented as achieving subgoals, some of which achieved simultaneously and independently (AND)
- Up to now, only dealt with OR options




AND/OR search

- We must examine several nodes simultaneously when choosing the next move













An Example

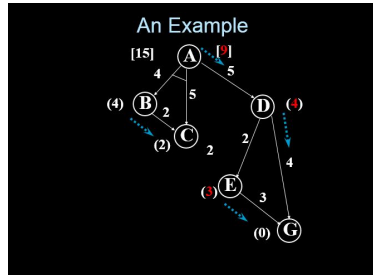

(8) A

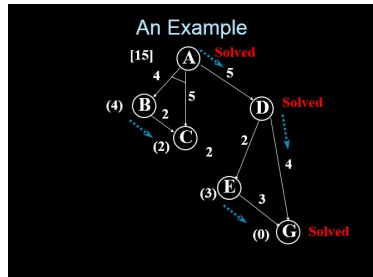













AO* Search

- 1 Place the starting node s on open.
- 2 Using the search tree constructed thus far, compute the most promising solution tree T_0 .
- 3 Select a node n that is both on open and a part of T_0 . Remove n from open and place it on closed.
- 4 If n is a terminal goal node, label n as solved. If the solution of n results in any of n 's ancestors being solved, label all the ancestors as solved. If the start node s is solved, exit with success where T_0 is the solution tree. Remove from open all nodes with a solved ancestor.
- 5 If n is not a solvable node (operators cannot be applied), label n as unsolvable. If the start node is labeled as unsolvable, exit with failure. If any of n 's ancestors become unsolvable because n is, label them unsolvable as well. Remove from open all nodes with unsolvable ancestors.



AO* Search

- 6 Otherwise, expand node n generating all of its successors. For each such successor node that contains more than one subproblem, generate their successors to give individual subproblems. Attach to each newly generated node a back pointer to its predecessor. Compute the cost estimate h^* for each newly generated node and place all such nodes that do not yet have descendants on open. Next, recompute the values of h^* at n and each ancestor of n .
- 7 Return to step 2.



Thank You!

